


Open Source MLOps: Course Intro

Alex Kim @alex000kim 

Course Instructor: Alex Kim




- Physics -> Software Engineering -> Data Science -> MLOps
- Currently, MLOps Engineer at Iterative.ai
- Worked on ML projects involving satellite image and video data, log analytics, mining sensor data, etc.
- Developed and taught DS & ML curriculums for EdTech bootcamps and universities
- Some facts about me:
 - Live in Montreal, 
 - Dad to two toddlers
 - Enjoy baking
- @alex000kim on [GitHub](#), [Twitter](#) and [LinkedIn](#)

What's MLOps?


Best practices, standards and tools for ML application development and deployment

- Dataset Management & Versioning
- Model Development
- ML Experiment Tracking
- CI/CD for model re-training
- Model Monitoring
- Team collaboration
- Progress reporting
- ...



About the Course

- Focus on “forest behind the trees”
- Real end-to-end project
- Detailed step-by-step instructions




P. KLAMMER

Open-source MLOps

When in doubt, go with Open-Source Software:

- Community and support
- Low/zero cost
- Public roadmap and discussions




Week 1: Kick-Starting an ML Project

- Apply best practices for establishing ML project structure and dependencies management
- Manage project dependencies with pip and virtualenv
- Version datasets with DVC



Week 2: ML Pipelines, Reproducibility and Experimentation


- Refactor a Jupyter notebook into a reproducible ML pipeline
- Version artifacts of an ML pipeline in a remote storage
- Iterate over a large number of ML experiments in a disciplined way



Week 3:

CI/CD for ML and ML-based Web API

- Learn the basics of CI/CD
- Leverage the power of CI/CD tools for ML projects with CML
- Integrate an ML model into the FastAPI framework
- Build and test a Docker container running a web API service
- Deploy the resulting Docker container to cloud



Week 4: Data Drift Monitoring for ML Projects

- Distinguish between application monitoring and ML monitoring
- Use Alibi Detect framework to detect data drift



What's your
experience
with ML and
MLOps?

